



COMP 206 – Introduction to Software Systems

Lecture 13 – Working with BMP Images


October 19, 2018

We are ready for (really useful) images!





Human Vision (not testable)

- How do people see?
 - One of the fundamental questions in early science. Major player: [Hermann von Helmholtz](#) 1821 – 1894. He understood that the eye has a lens which projects light using physical laws. What happens afterwards?




Human Vision (not testable)

- How do people see?
 - Now we know about the retina with light-sensitive rods and cones in precise pattern.
 - Optical nerve transmits sensed light to visual cortex arranged in "retino-topic" fashion (like a grid) for several layers. Nobel prize for [David H. Hubel](#) and [Torsten N. Wiesel](#) in 1981.




How do cameras form images?

- Project light with a lens that has similar behavior to our eye.
- Instead of a retina, light sensitive electronics (CCD or CMOS), count arriving photons. Each is tuned for a color, we call **R**ed, **G**reen, or **B**lue. The **R****G****B** values at one spot are called a *pixel*.
- Each pixel is read off as 3 integer values (binary memory!)



How to store images as a file on disk?

- An image is a 2D grid of pixels:
 - num_rows = height
 - num_cols = width
 - num_colors: how many per pixel could be 3 for RGB, 1 for b/w, or 4 for RGBA (alpha = transparency)
 - Bits per pixel: what size of integer is needed to store each?
- 2 additional types of data:
 - A header, holds information fields such as the image size, compression, color depth
 - Padding, almost always present to align the elements into 4 or 8 byte boundaries (details coming)

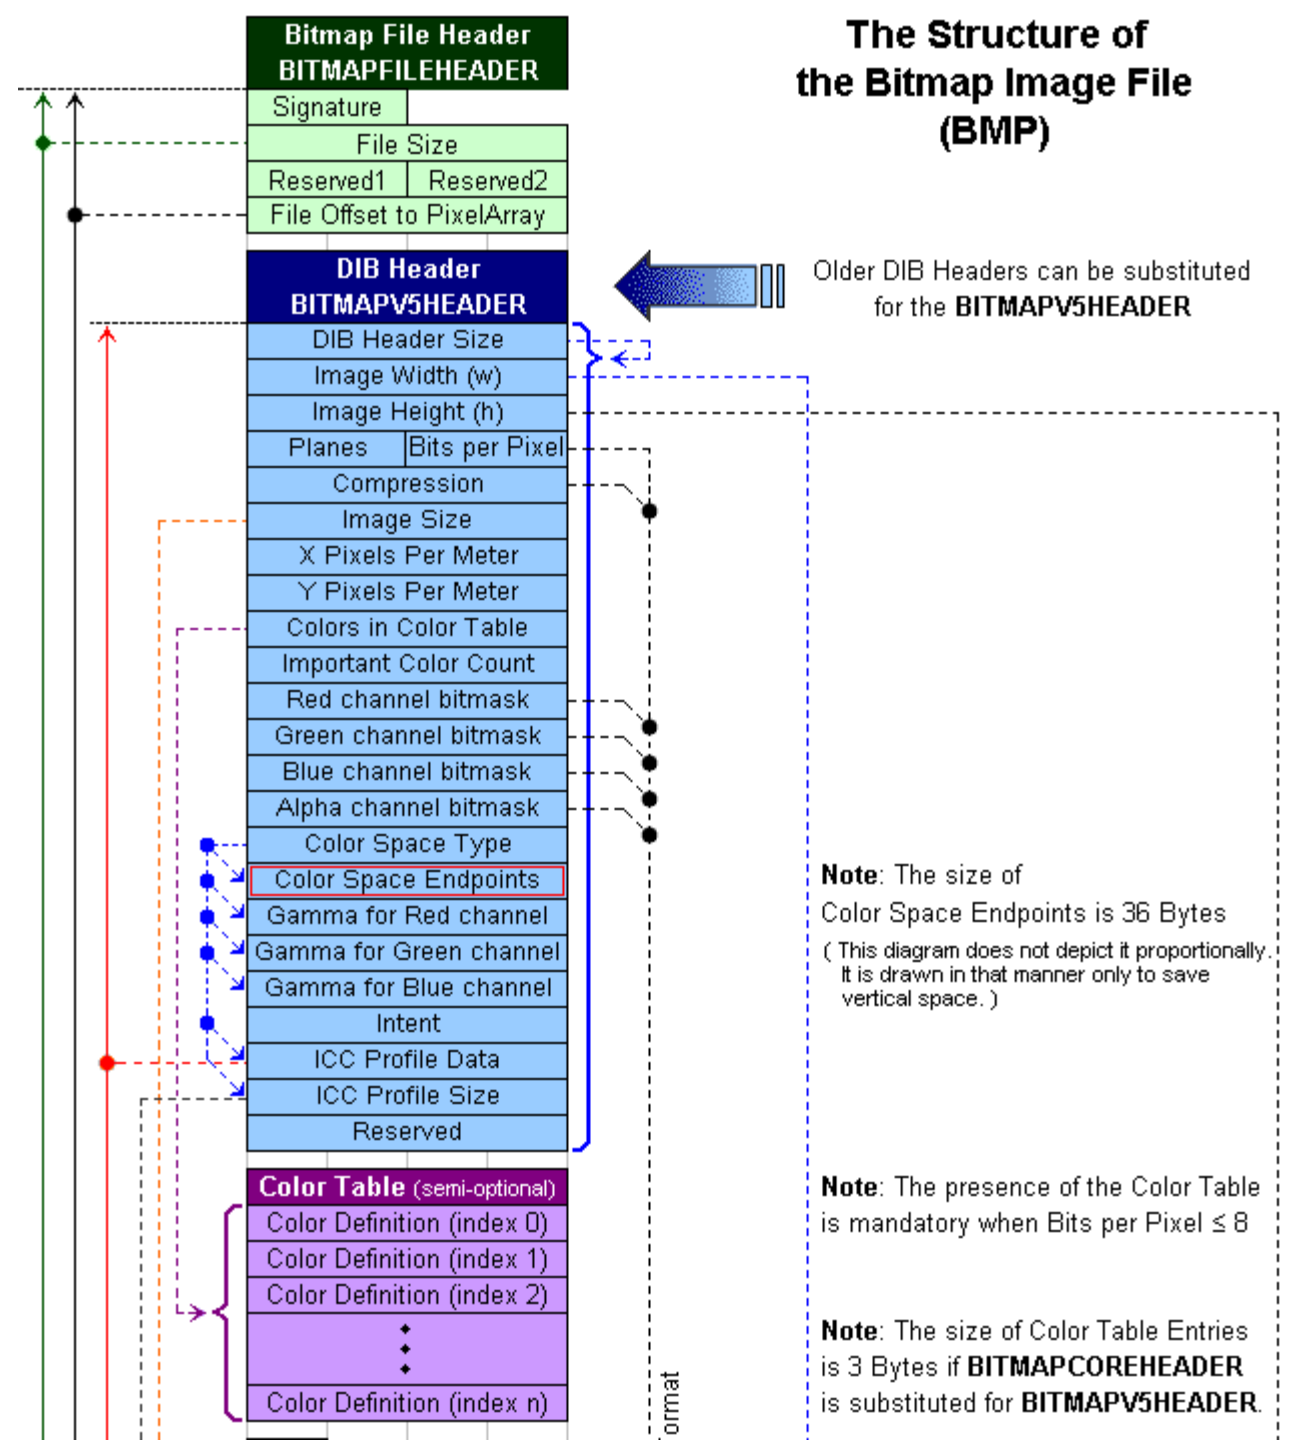


Many image file formats

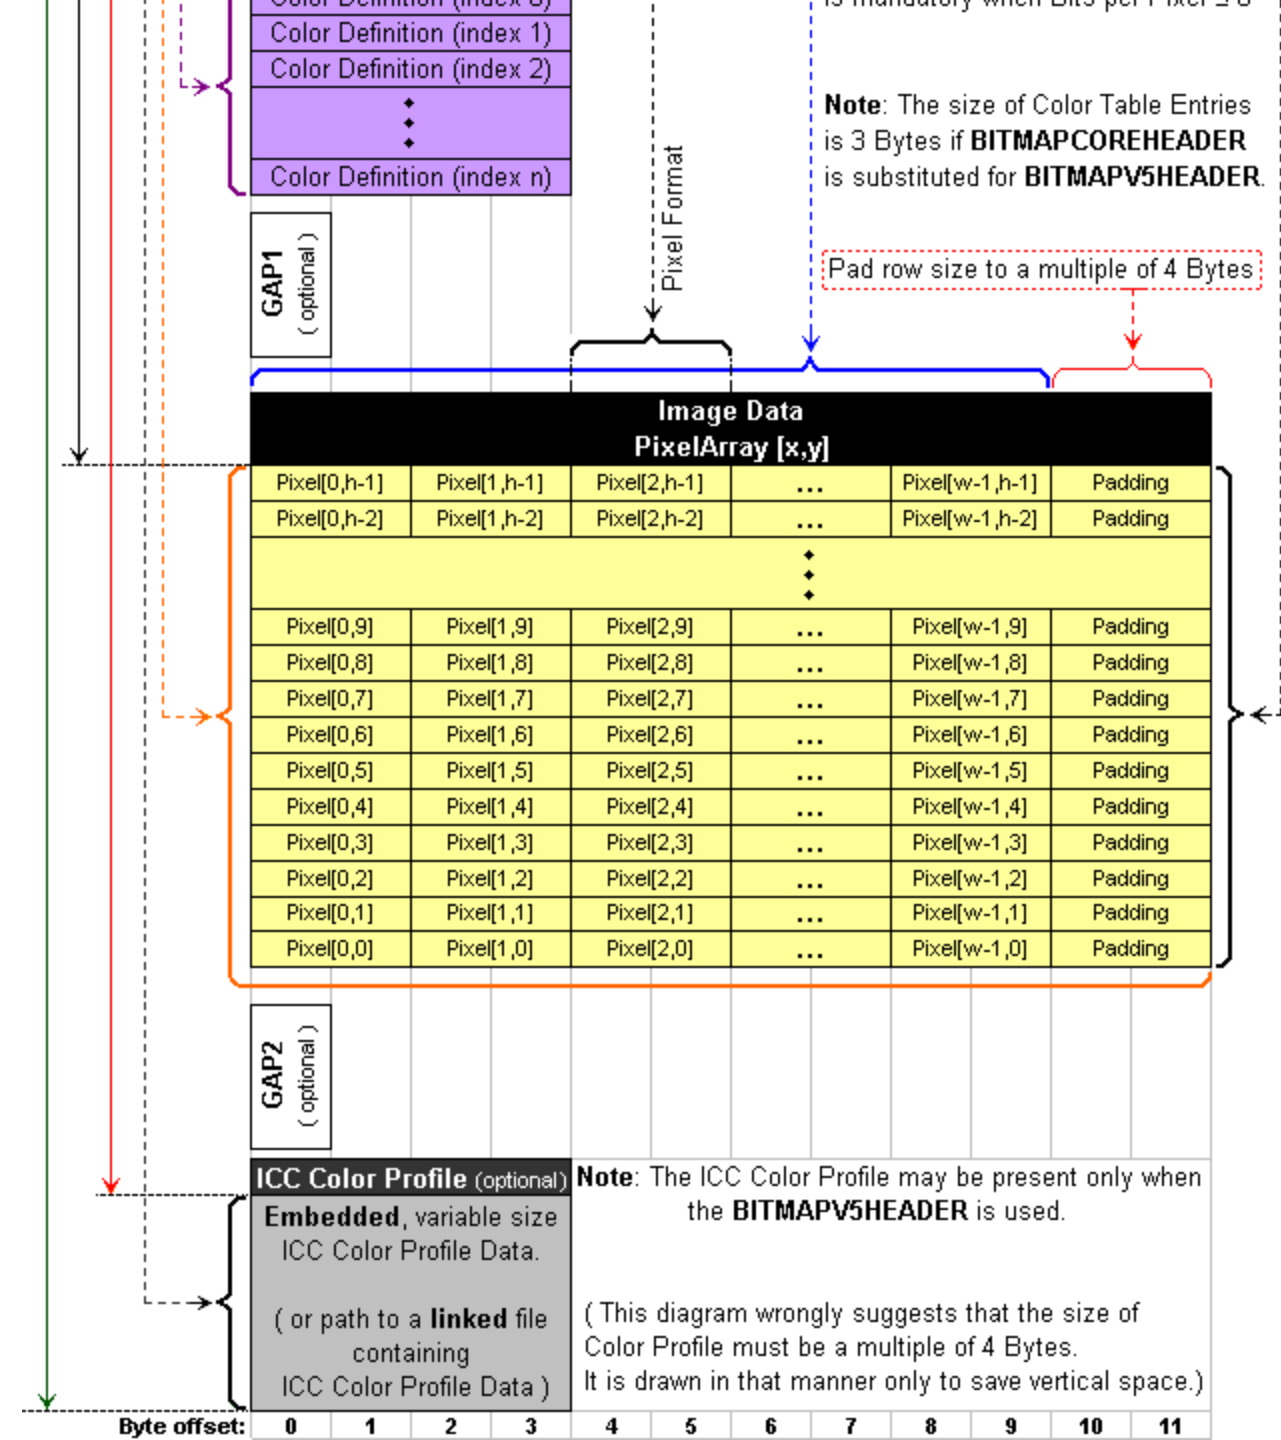
- BMP: "Windows bitmap"
- JPG: Joint Photographic Experts Group
- PNG: Portable Network Graphics
- TIFF: Tagged Image File Format
- GIF: Graphics Interchange Format
- Of these, only BMP is testable in 206 and will be used in A3.



Bitmap File (BMP) Example first half



Bitmap File (BMP) Example second half



How can we read a BMP file using C?

- What works well:
 - Check the magic number:
 - If it matches very likely it follows the rules
 - File size field: makes it easy to access all of the data
 - Width and height, allows finding a specific pixel
 - Opening with code like “rb”
- What we must avoid:
 - Checking for ASCII code values: space, newline, etc
 - Attempting to use “atoi” “atof”, these are “ascii to ...”
 - If we open with “r” alone (no b), C will do some of this automatically and cause us problems.
 - fgets, fscanf also typically bad choices, mean to work with text

Example

- Github:
ExampleCode/
Lecture13 folder
- Note "18" is the byte for width using the chart above
- Ensure you understand how to read the chart (needed for A3)

```
int main(){

    // Open a binary bmp file
    FILE *bmpfile = fopen( "utah.bmp", "rb" );

    if( bmpfile == NULL ){
        printf( "I was unable to open the file utah.bmp.\n" );
        return -1;
    }

    // Read the B and M characters into chars
    char b, m;
    fread ( &b, 1, 1, bmpfile );
    fread ( &m, 1, 1, bmpfile );

    // Print the B and M to terminal
    printf( "The first byte was: %c.\n", b );
    printf( "The second byte was: %c.\n", m );

    // Read the overall file size
    unsigned int overallFileSize;
    fread( &overallFileSize, 1, sizeof(unsigned int), bmpfile );
    printf( "The size was: %d.\n", overallFileSize );

    // Rewind file pointer to the beginning and read the entire contents.
    rewind(bmpfile);

    char imageData[overallFileSize];
    if( fread( imageData, 1, overallFileSize, bmpfile ) != overallFileSize ){
        printf( "I was unable to read the requested %d bytes.\n", overallFileSize );
        return -1;
    }

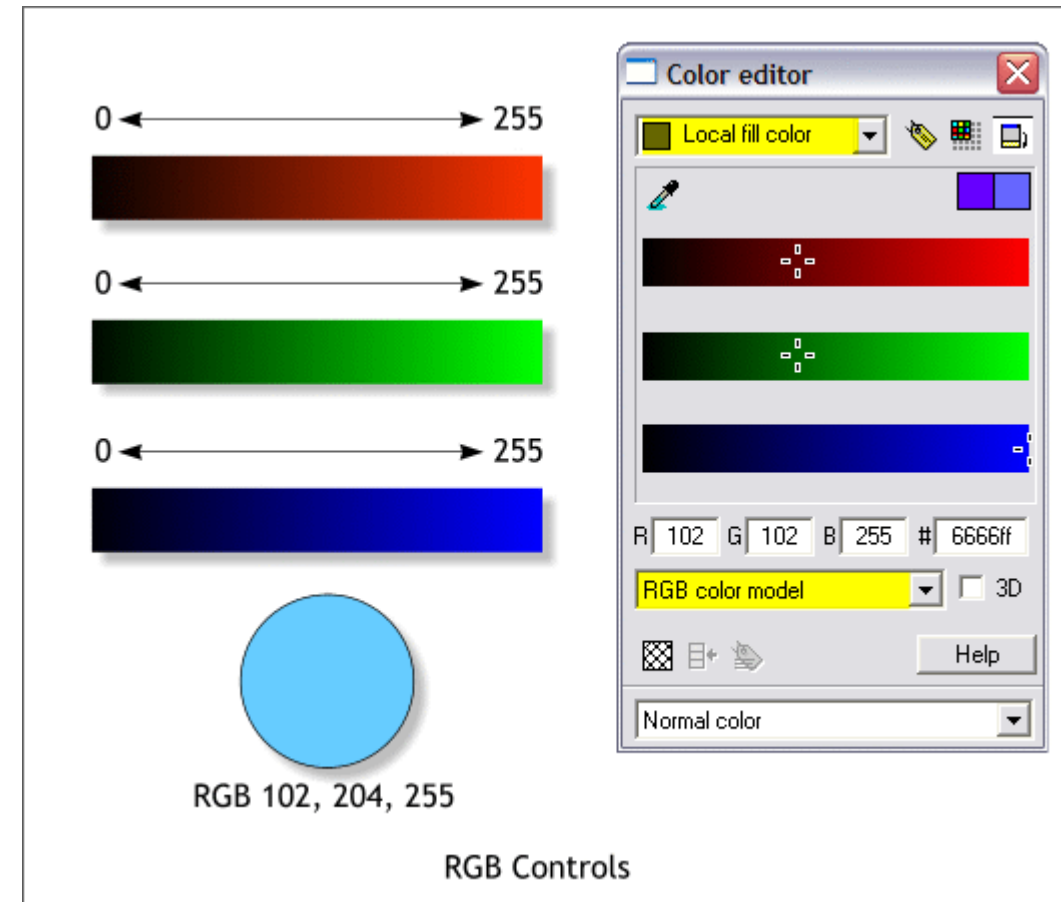
    // Read the width size into unsigned int (hope = 500 since this is the width of utah.bmp)
    unsigned int* wp = (unsigned int*)(imageData+18);
    unsigned int width = *wp;

    // Print the width size to terminal
    printf( "The width is: %d.\n", width );

    return 0;
}
```

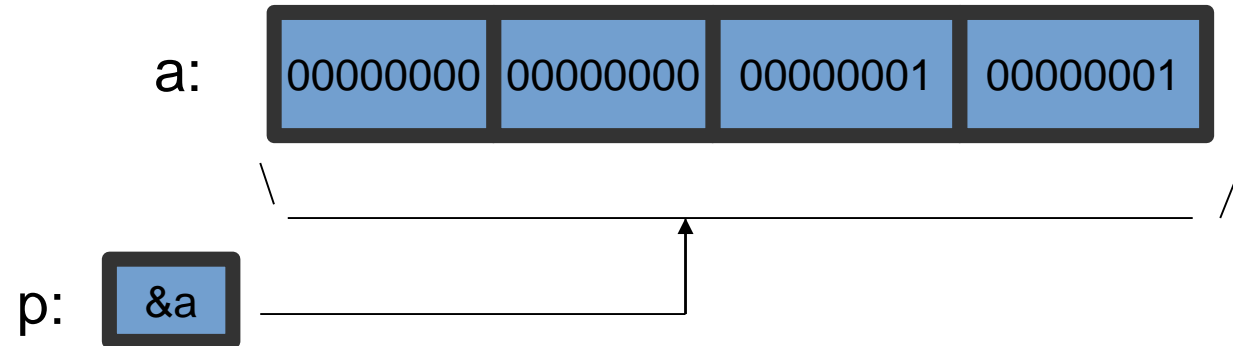
Now that we have the data...

- Each color of each pixel is stored as an integer between 0 and 255 (one byte):
 - Easiest way to work with these in C: represent each as an unsigned char
- Other items such as the length and width are 4 byte integers
- It's time to get brave and learn some new tools to work with mixed memory!



Normal pointer use

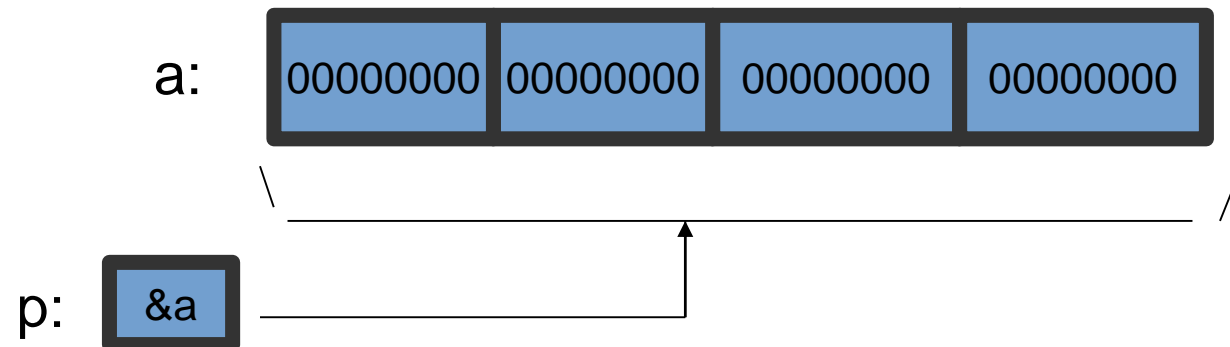
- `int *p;` -> a pointer to an integer value
- `int a = 257;` -> an integer variable
- `p = &a;` -> p now holds the **address of a**
(points to a)



All is well here. But what if we change the types?

Pointers can re-interpret memory

- `int *p;` -> a pointer to an integer value
- `char a[4];` -> a character array (string)
- `memset(a, '\0', 4);`
- `p = (int*)a;` -> p now holds the **address of a** (points to a)



- This is OK, and treats the whole string as integer data (equivalent integer value 0)


Example Code

- `ExampleCode/Lecture13-ImageFiles/pointer_conversion.c`
- Try this out, think about changing values in the C string – what happens to the pointer? This can all seem weird, so it's important you really try it yourself!

What about the other direction?

- `char *p;` -> a pointer to a character value
- `int a = 257;` -> an integer variable
- `p = (char*)&a;` -> p now holds the **address of a**,

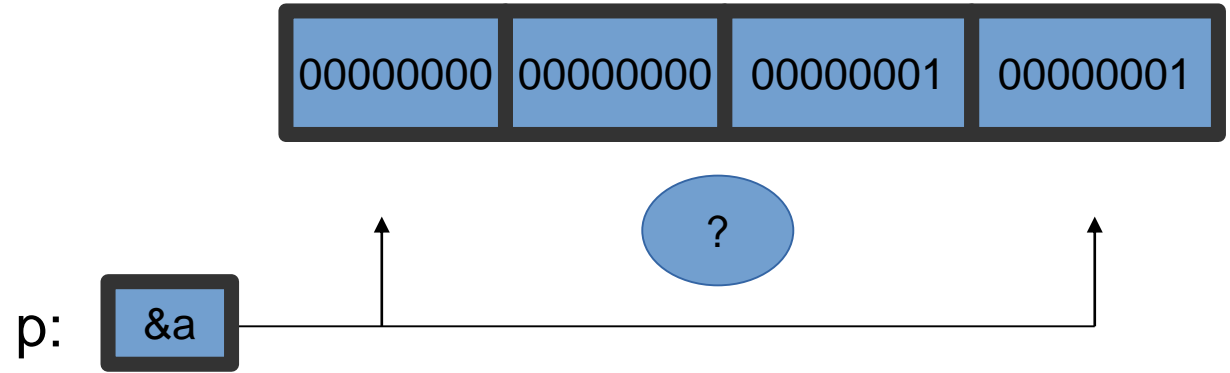
but will interpret it as character memory, 1 byte only, when dereferenced!



- Think carefully about what character byte `p` will address. What will the value of `*p` be?

Pointers and Type Conversion



```
int a = 257;  
char *p = (char*)&a;  
printf( "The value at *p is %d.\n", *p );
```



The output of this program depends on “endianess”. Sample code GitHub:
[endianness_test.c](#)

Word Endianness

- The order that bytes within an integer are stored in memory is a convention, named Endianness, and there is no right answer.
- Most systems we deal with will be Little-endian, but there are major exceptions (the Sun company)
- It is always better to check than to assume




Endianess intuition helper

- Little or big? Name determined by the "significance" of the byte at the first address:
 - Little: the least significant comes first
 - Big: the most significant comes first
- Humans always write numbers in Big Endian, why would most computers use Little?
 - Think about doing addition with carry-over
 - We process right to left and "carry"
 - Computer add a lot, and are most efficient accessing memory "in order"



00000000	00000000	00000000	11111111
00000000	00000000	00000000	00000001

Intuition Example






256



Computed using
Big-Endian



Computed using
Little-Endian



Intuition Example

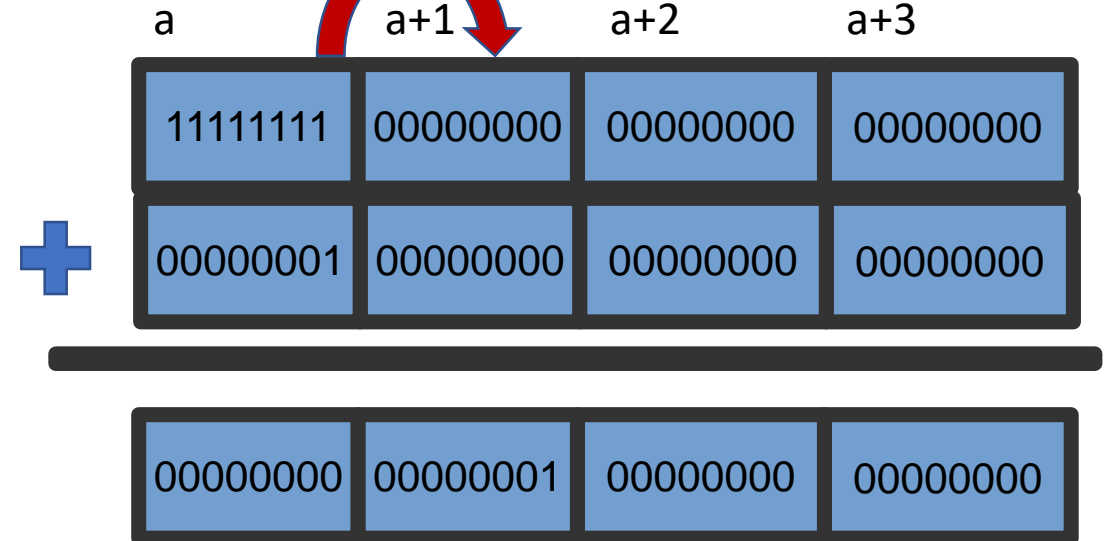



Computed using
Big-Endian

Computed using
Little-Endian

“Carry over”
backward

“Carry over”
forward



Endianness and Images

- We began to think about this topic because we started to convert between memory of different lengths. Will it really matter for images?
- Not for BMPs, probably. BMP files are explicitly specified to use Little Endian, which is the format of most machines (all x86 compatible):
 - To tell on your machine, try running the sample code or type “uname -a” on the terminal. If you see “x86” anywhere, it’s Little-Endian
- So why does this matter? Because we are not so lucky in many cases:
 - PNG images: <https://www.w3.org/TR/PNG/#3byteOrder>
 - The Internet mostly uses Big Endian! Hence B.E also called “Network Byte Order”

Reading and Writing Pixels?

- It's time to get going and change some data. Let's set all the pixels to white!
 - bmp_2_white.c in ExampleCode

```
// To modify the pixel data, we need to make sure we dont
// mess up the image header. Otherwise, image viewers will
// no longer understand the file. Let's read the offset
// and move a pointer forward.
unsigned int* offsetp = (unsigned int*)(imageData+10);
unsigned int offset = *offsetp;
char *pixel_data = imageData + offset;

// Now, indexing to pixel_data will only change the "visible"
// data in the image. White is all of R, G and B = 255.
// Warning! This changes the "padding" parts of the image also.
// For A3, you will need a more sophisticated way of managing
// the image data.
for( int pixel=0; pixel< overallFileSize-offset; pixel++ )
    pixel_data[pixel] = 255;

// Time to output, just dump the binary data, inverse of reading.
FILE* out_file = fopen( "utah_white.bmp", "wb" );
if( out_file == NULL ){
    printf( "Unable to open utah_white.bmp for writing.\n" );
    return -1;
}
fwrite( imageData, 1, overallFileSize, out_file );
```

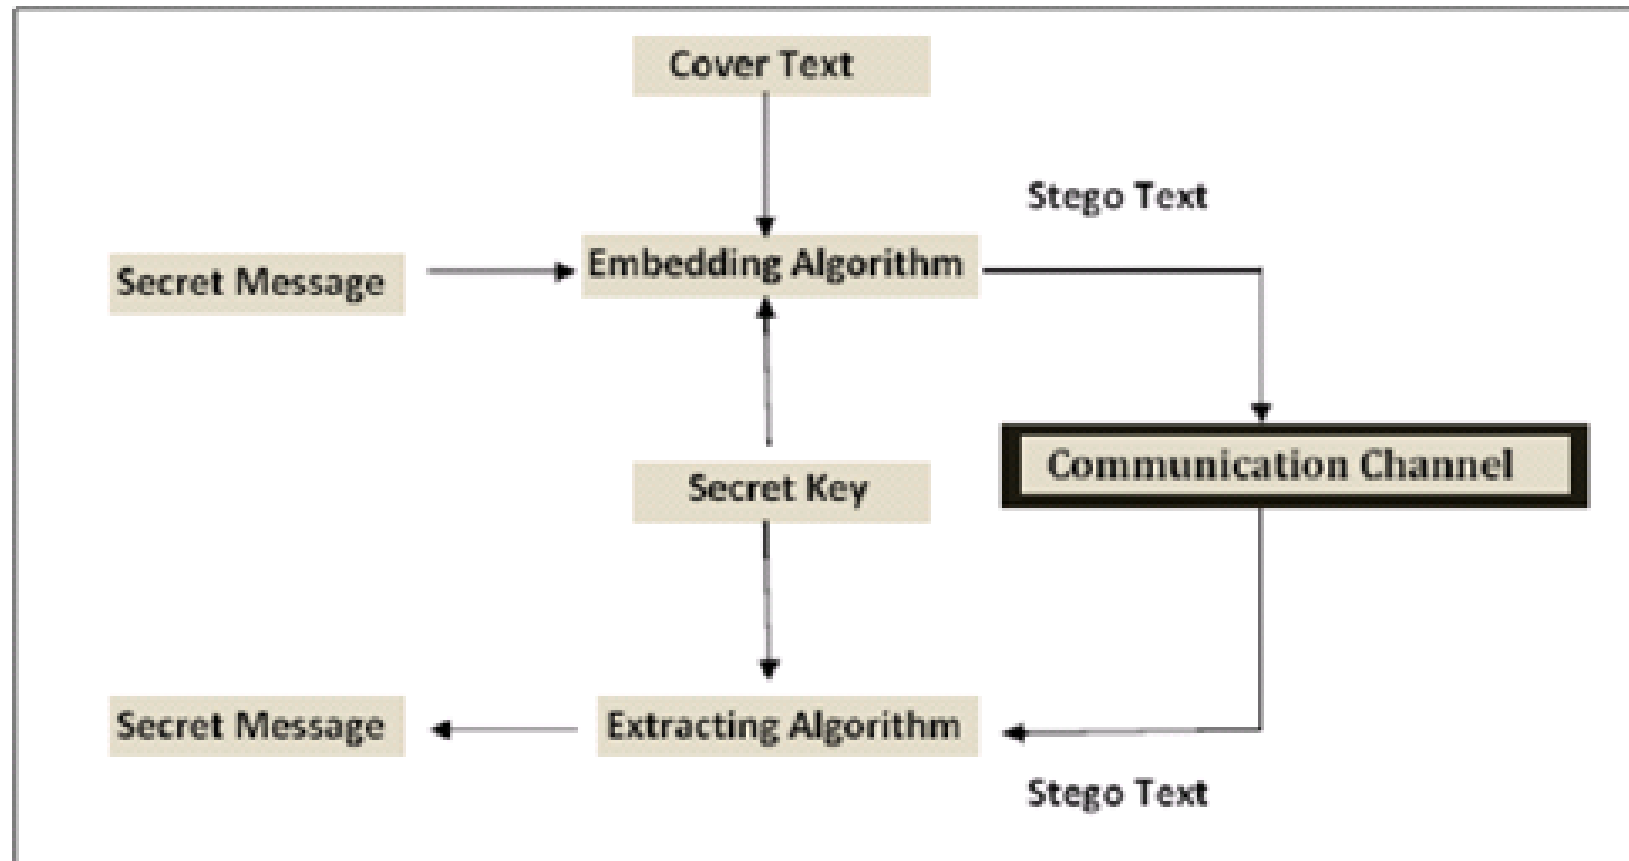

A motivating example: steganography



Steganography

- ...is concerned with concealing the fact that a secret message is being sent as well as concealing the contents of the message.
- Whereas cryptography is the practice of protecting the contents of a message alone, steganography is concerned with concealing the fact that a secret message is being sent as well as concealing the contents of the message.

Diagram of a Steganography Software System




Live Demo!

steg_encoder.c and steg_decoder.c

- Two files posted in this `ExampleCode/Lecture13` folder on GitHub
- We have not yet covered everything needed to understand them, but we will demo quickly to show what we're after
- The rest of the details in the last part of these slides

Digital Steganography

LSB IN IMAGES



144	141	81
-----	-----	----

10010000 10001101 01010001

Hidden message: 101001...

145	140	81
-----	-----	----

1001000**1** 1000110**0** 0101000**1**

146	142	81
-----	-----	----

100100**10** 100011**10** 010100**01**

Elements Required for In-Image Steganography

- Read and write binary image data (DONE!)
- Read and change pixel values (DONE!)
- Steganography requires us to work with **bits** directly, which we haven't considered yet:
 - View the text string in binary form, so we can access one bit at a time
 - Ability to modify only a single bit (the LSB) of each pixel
 - Ability to extract the LSB again for decoding

Bit-wise Operations

Shifts:


- `bit_arg << shift_arg`
 - Shifts bits to of `bit_arg` `shift_arg` places to the left -- equivalent to multiplication by $2^{\text{shift_arg}}$
- `bit_arg >> shift_arg`
 - Shifts bits to of `bit_arg` `shift_arg` places to the right -- equivalent to integer division by $2^{\text{shift_arg}}$

Bit-wise logic:

- `left_arg & right_arg`
 - Takes the bitwise AND of `left_arg` and `right_arg`
- `left_arg | right_arg`
 - Takes the bitwise OR of `left_arg` and `right_arg`
- `left_arg ^ right_arg`
 - Takes the bitwise XOR of `left_arg` and `right_arg` (one or the other but not both)
- `~arg`
 - Takes the bitwise complement of `arg`


Bit-shift Operators

- Moves the existing bits a specified number of positions left or right.
 - When a bit hits the edge, it is lost
 - New bits always take 0 (for unsigned – we do not cover signed bit shifting in 206)
- Note, shifting is similar to multiplying or dividing by powers of 2



Multi-byte shifting

- Treats the true integer value. That is, we do not have to think about address ordering here. If a bit hits the boundary of its byte, it seamlessly moves to the next using significance order.
- Now, only the least and most significant byte boundaries cause “loss”



Bitwise Logical Operators

- Each applies a *truth table* to the bits in its arguments, one at a time
- Logical AND and logical OR truth tables:

	0	0	1	1		0	0	1	1
&	0	1	0	1		0	1	0	1
<hr/>						<hr/>			
	0	0	0	1		0	1	1	1

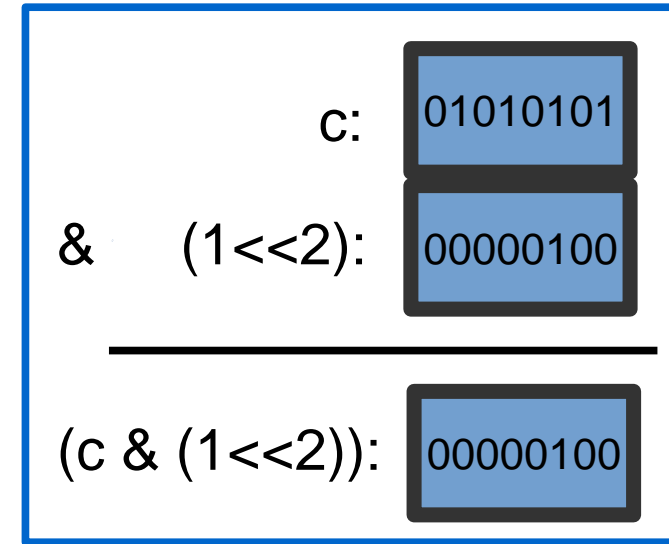
- When you apply `a & b`, these operations are applied to all of the bits in `a` and `b`, 1 bit at a time

Integrated Bitwise Example

- Check the value of bit 3:

```
char c = 85;  
if( (c & (1<<2)) > 0 )  
    printf( "bit 3 of c is a 1!\n" );  
else  
    printf( "bit 3 of c is a 0!\n" );
```

- Larger example posted to Github:
 - [bit_reporting.c](#)



More Bitwise Examples

- Set the value of bit 8 to 1
- Count the number of 1's
- Find the first 0 starting from the right

Back to Steganography

- Recall: if we encode a string within the “lowest-order” bits of an image, it is nearly invisible
- We will examine the example code, posted on Github for this lecture
 - [steg_encoder.c](#)
 - [steg_decoder.c](#)



Steganography

Learning Outcomes

- Understand how our code interacts with external systems:
 - BMP file format: defines how our data is structured, forces our code to ignore the header, binary data format
 - Image viewer: displays our data, determines which parts of the data matter
 - Our program is called on the command-line, can interact with BASH functions.
 - **CHALLENGE #1:** How would you “hide” the contents of another file inside the image?
- MOSTLY COMPLETE!**

Steganography

Learning Outcomes

- Understand how the two pieces of our code interact with each other:
 - **CHALLENGE #2:** What if we wanted to change the bit used to hide our data?
 - What if the decoder wanted to switch to JPG data?
 - How does the decoder know the length of the message?
 - **CHALLENGE #3:** Could we make each side more robust to ensure the image and message are “proper”?

Read More

- K & R textbook Section 2.9